Metamathematics of Algorithmic Composition

Michael Gogins

michael.gogins@gmail.com
https://michaelgogins.tumblr.com

May 2025

This essay recounts my journey towards a deeper understanding of the mathematical foundations and computational complexity of algorithmic music composition. I do discuss some specific algorithms used by myself and other composers, but my primary focus is fundamental limits and possibilities of algorithmic composition, by analogy with metalogic, metamathematics, and computability theory. I discuss implications from these foundations for the practice and future of algorithmic composition.

Introduction

In 1983, when I was a returning undergraduate in comparative religion at the University of Washington, I attended a lecture on fractals by Benoit Mandelbrot, discoverer of the set named after him (Mandelbrot 1982; Peitgen et al. 2004). Briefly, given the quadratic recurrence equation $z_{n+1} = z_n^2 + c$, the Mandelbrot set consists of all points c in the complex plane for which z, starting with $z = 0$, does not approach infinity as the equation is iterated. Then, for each point c in the Mandelbrot set there is a connected Julia set, consisting of all points z for which z_n does not approach infinity as the equation is iterated. Mandelbrot showed slides illustrating how any point in the Mandelbrot set can be used as the generating parameter of a Julia set, and how a plot of the neighborhood near a point in the Mandelbrot set closely resembles the plot of the corresponding Julia set (Lei 1990) (this has recently been proved (Kawahira and Kisaka 2018)). In short, the Mandelbrot is a *parametric map* of all Julia sets. There is now an extensive and still growing literature on the Mandelbrot set and Julia sets.

Already at the time of Mandelbrot's lecture, I was developing an interest in computer music and algorithmic composition, in particular, algorithmic composition based on fractals. During the lecture it occurred to me that if I zoomed into a plot of the

Mandelbrot set, searching for interesting-looking regions, I could plot the Julia set for a point in that region, and I could then somehow translate that Julia set into a musical score (Gogins 1992a). I have worked out several variations of this idea.

Mandelbrot Set/Julia Set

The composer explores the actual Mandelbrot set. When a region seems interesting, the composer selects a point in it, and the corresponding Julia set is first generated, and then translated to a musical score. The 2-dimensional plot of the Julia set is mapped more or less directly onto a 3-dimensional piano-roll type score, with the x axis representing time, the y axis representing pitch, and the color of a point in the plot representing choice of instrument.

Parametric Map of Iterated Function Systems

An iterated function system (IFS) is a Hutchinson operator, a set of contractive affine transformations. Iterating the operator upon any initial set takes that set to a fixed point that is a fractal (Barnsley and Demko 1985; Barnsley, 1993 [1988][a]). The Collage Theorem proves that an IFS can approximate any set as closely as desired (Barnsley and Sloan 1989; Barnsley, 1993 [1988][a]). It follows that this method is *compositionally universal* (Gogins 1992a, In preparation), that is, capable of directly generating any finite score. It is also possible to generate a parametric map of any subset of IFSs, but because the Hutchinson operators have many more than two parameters, generating a parametric map (an analogue of the plot of the Mandelbrot set in the complex plane) for dozens or hundreds of parameters requires the use of a Hilbert index, which maps points in an N-dimensional space, such as a plane, cube, or hypercube, to a 1- or 2-dimensional sequence of numbers (Patrick, Anderson, and Bechtel 1968). The index is constructed in such a way that neighboring points in the N-dimensional space have nearby indices in the 1- or 2-dimensional sequence. The key idea is to recursively subdivide the N-dimensional space into smaller planes, cubes, or hypercubes, called cells. Each cell is assigned a unique index based on its position within the overall space. Subdivision continues recursively until a desired level of detail is reached. To determine the Hilbert index of a specific point in the N-dimensional space, start with the largest cell that contains the point, level $j = 0$. Then, working in arithmetic to the base N, subdivide that cell into N sub-cells for level 1, and select the sub-cell that contains the point. If it is the kth sub-cell at level j, then add $((k + 1)/N)^{-j}$ to the index. Repeat this process recursively until the smallest cell containing the point is reached. The index of that cell is the Hilbert index for the point. Hilbert indices work because all metric spaces have the same cardinality; hence there is always a one-to-one mapping between points in any N-dimensional space and points on the line or, more usefully, the plane. In any event, the plot of the IFS is translated to a score in much the same way as for the plot of a Julia set.

    I experimented with both methods, doing parametric composition by zooming in on interesting regions of the map, generating and rendering scores, exploring points near

scores that seemed promising, and iterating this process. However, scores generated from Julia sets have too much of a sameness, and producing a parametric map of more than just a few points for IFSs simply takes way too much time. In other words, computing parametric maps for IFSs is *computationally intractable*.

As I pursued algorithmic composition, I found that this question of intractability is not merely a practical problem, but has profound mathematical and philosophical roots, beginning with Pythagoras and continuing on through the hierarchy of complexity classes in theoretical computer science.

Philosophical Context

Music has since Pythagoras (C. Huffman 2024; C. A. Huffman 2014) been understood by some as an intellectual paradigm and to reveal, through harmony that is both numerical and audible, the structure of reality. For this reason music was a central part of the *quadrivium*, the standard curriculum of liberal arts in Western higher education from late Antiquity through the Renaissance.

The project of understanding reality through number advanced from Pythagoras, through Leibniz' hope for a *characteristica universalis*, a symbolic language that could express all rational thought (Davis 2018), to the logicism of Russell (Tennant 2023), Hilbert (Zach 2023), and others, that sought to derive all mathematical truths from formal logic. In one of the major achievements of all philosophy, Kurt Gödel (Gödel 1986) demonstrated in his incompleteness theorems that logicism cannot be fully implemented, because there exist true statements of logic that cannot be proved. About five years later, Alan Turing proved the Halting Theorem: it is impossible for any computer program to decide whether another, arbitrary computer program will halt. Researchers are still exploring the consequences of these theorems.

One major result has been the elucidation of a provable hierarchy of complexity classes for problems that are solvable by computer (Arora and Barak 2009). Other important results are proofs there are completely abstract machines that are not incomplete; these are called super-Turing and depend, one way or another, on doing arithmetic with real numbers (almost all real numbers are not Turing computable, because their decimal representations never terminate).

Complexity Classes

The complexity classes are based on the capabilities and limitations of Turing machines. A Turing machine is an abstract, idealized computer with infinite memory that can run for an infinite number of steps. It may halt after a finite number of steps and produce a result, or never halt. A universal Turing machine is one that can emulate any other Turing machine. The widely accepted Church-Turing thesis holds that *anything* computable by a definite, step by step procedure is computable by a universal Turing machine (Copeland 2024). Our contemporary digital computers are universal Turing machines – or would be, if they had infinite memory and infinite time to run.

For our purposes, the important complexity classes, assuming that complexity is defined as run time on a universal Turing machine, are as follows.

Super-Turing (ST) Problems that have *mathematical* solutions, but the solutions are not *Turing computable.*

Recursively enumerable (RE) Also called semi-decidable. Problems where a solution can be verified by a Turing machine, but proving that a solution does not exist is not always possible. If a solution does exist, then a Turing machine can get as close as one likes to solving the problem.

Recursively computable (R) This is the same as Turing computable, or decidable. The problem can be decided in a finite number of steps; that is, the problem can be either solved, or proved unsolvable.

Non-deterministic polynomial time (NP) Recursively computable, but the the time to solve increases faster than any polynomial function of the size of the problem. However, a solution can be *verified* in polynomial time.

Deterministic polynomial time (P) Recursively computable, and the time to *solve* is bounded by some polynomial function of the size of the problem.

The complexity classes, including ST, mirror central issues in philosophy and science. It is a primary open question of philosophy whether Nature herself is super-Turing. If so, then human thought, including musical composition, might as part of Nature also be super-Turing. If not, then human thought is at most RE and could be emulated as closely as one likes by a Turing machine.

Yet in either case, scientific theories are RE or less, because it must be possible to compare the predictions of the theory to observations of Nature that are finite in number and of only finite precision. Nobody can see how Nature or human thought being super-Turing could ever be proved, so it is often held that any physical process can be effectively emulated by a Turing machine; this is the *physical* Church-Turing thesis (Aaronson 2005; Copeland 2024).

The closely related question whether NP is contained in P is one of the most important open questions in science. Most mathematicians and scientists believe, for a variety of overlapping reasons, that $\mathsf{P} \neq \mathsf{NP}$.

If $\mathsf{P} = \mathsf{NP}$ can be proved, it then becomes possible in principle to automatically solve all problems of a size that human beings can grasp. And this means that, given the problem of composing a certain piece of music to specification, a composer is not needed. Algorithmic composition alone can do the job. In terms of parametric composition, computing a finite subset of the parametric map for IFSs would be in P (with some limitations).

But if $\mathsf{P} \neq \mathsf{NP}$ can be proved, then we will know that there are problems that computers simply cannot solve, but human beings *perhaps* can solve. And this means that, given the problem of composing a certain piece of music to specification, a composer is a good thing to have around. Not only that, but the use of algorithmic composition does not

change the benefits of involving a composer. In terms of parametric composition, that is not because of the indexing, but because of the time required to evaluate a *piece* for each point in the parametric map. Recall that computing a piece as a Julia set may be in NP; therefore, computing a parametric map of pieces that includes such a piece would also be in NP. However, looking up the set of parameters for a given point in the map is in P. It may be only $O(log\ n)$ or, for multi-dimensional searches, $O(n^c), 0 < c < 1$.

If neither can be proved, then we will just never know.

Algorithmic Composition in Context

It is now possible to return to the analysis of algorithmic composition with more understanding. I use parametric composition using Julia sets as standing in for algorithmic composition in general. This is justified by the fact that just like Julia sets, many other compositional algorithms use real numbers or depend in some way on chaotic behavior.

The starting point here is Braverman's proof (Braverman and Yampolsky 2006, 2009) that most (hyperbolic) Julia sets are computable in P (measure 1); some (parabolic) Julia sets are computable, but in NP or worse (measure 0); and some with Seigel discs or Cremer points (neither hyperbolic nor parabolic) are *uncomputable* (also measure 0). Hence it turns out that parametric composition spans the entire hierarchy of complexity classes.

ST Julia sets with Seigel discs or Cremer points may be of musical interest in the abstract, but they are uncomputable. Picking Julia set parameters at random is not at all likely to find one of these.

RE The Mandelbrot set, properly speaking, is not *recursively computable*, *i.e.* not Turing computable (Blum and Smale 1993). The plots we make of the set are approximations. Hertling showed that although the Mandelbrot set is not recursively computable, it is nevertheless *recursively enumerable* (Hertling 2005); given enough time, one can approximate the actual set as closely as one likes.

NP If any parabolic Julia sets are musically interesting, then computing a parametric map that is complete enough to be useful for composers might remain forever computationally intractable. Picking Julia set parameters at random is not at all likely to find one of these. However, as Julia set parameters are chosen closer and closer to a parabolic point, the Julia sets get more and more difficult to compute and the appearance of computer generated plots (or mappings to scores) fluctuates wildly.

P But if P = NP then in principle it is possible to produce a useful parametric map in polynomial time: no set with Seigel discs or Cremer points, but all hyperbolic points and as close as we like to all parabolic points. In that case, parametric composition might open up a vast new world for composers.

To restate the question, if P = NP, then a high-resolution parametric map of considerable subsets of compositions would, as it were, display symmetries and patterns,

such as those found in the Mandelbrot set, that could assist in parametric composition. It would afford a God's-eye view of (almost) all possible structure in music, and would partly overcome the irreducibility of algorithmic composition. In other words, $\mathsf{P} = \mathsf{NP}$ could imply that algorithmic composition can be made *intelligible*.

But if $\mathsf{P} \neq \mathsf{NP}$ parametric maps must exclude not only all sets with Seigel discs and Cremer points, but also many parabolic points and larger neighborhoods of them. This would to some extent fracture and obstruct the intelligibility of algorithmic composition.

I would like to point out that, although almost all (nearly measure 1) of Julia sets are in P, not understanding those in NP or greater risks profound misunderstandings. To understand the actual mathematics requires continuous analysis and real numbers, whereas computing actual compositions can be done only with discrete methods and integers, and only on a physical computer. We must constantly remember that things may be mathematically true yet not computable, and other things may be computable in the abstract yet not computable in our physical world.

Another significant aspect of parametric composition in particular, and of algorithmic composition in general, is *computational irreducibility*. Although any Julia set approximately resembles the neighborhood of the Mandelbrot set near its generating parameter, almost every Julia set is the chaotic attractor of its generating equation. Therefore, the orbit of the Julia equation is *computationally irreducible* in the sense of Wolfram (Wolfram 1985), as proved by Zwirn (Zwirn 2015). The orbit of the equation cannot be determined by examining the equation, and it cannot except in trivial cases be determined even by mathematically analyzing the equation. In order to know the orbit, it is necessary to actually iterate the equation, that is, to run a program that computes the iterations. Even then, we can only obtain an approximation. What is true of Julia sets is true of IFSs or, for that matter, of any algorithm that is not just rudimentary. In order to know what an algorithm actually does, one must actually run the algorithm.

Before further exploring the mathematical foundations of algorithmic composition, I will provide some additional background relating to different software systems for algorithmic composition, and to machine learning, which can also be used to compose music.

Methods of Algorithmic Composition

To expand upon the concept, *algorithmic composition* is the use of computer software to write musical compositions that would be difficult or impossible to write without software. It does not include the use of notation software to write down music that one imagines (as that can be done with paper and pencil), nor does it include the use of audio workstation software to record and overdub music that one improvises (as that can be done with a tape recorder). In other words, algorithmic composition consists of all compositional techniques that are *idiomatic* to the computer. Of course, there is not just one method of algorithmic composition (Fernández and Vico 2013; Ariza 2023). A recent summary can be found in (McLean and Dean 2018). There is an obvious overlap with a more generic notion of *process music* or *generative music*, including Mozart's

musical dice game (humdrum.org 2023), the minimalism of Steve Reich (Reich 2023; Schwarz 1980) and Philip Glass (Potter 2002; Glass 2015), and the generative work of Brian Eno (Eno 2023). The commonality between algorithmic composition and process music is precisely the simplicity and clarity of the means versus the complexity and unpredictability of the results; in other words, yet again, irreducibility.

Here I should clarify this idea of irreducibility. It is not a binary choice, it is a spectrum. The minimum of irreducibility occurs when the composer simply writes down what he or she hears in his or her imagination. The maximum occurs when the composition is generated in an entirely random way, so that there is absolutely no way for the composer to predict, better than chance, any particular note or sequence of notes; but even then, there is a degree of musical intelligibility in that the texture of one random variable (e.g. white noise) can easily be distinguished from the texture generated by another random variable (e.g. brown noise). In the middle of the spectrum is an area where the composer does have some degree of insight into the kind of music that will be generated, even though the details cannot be predicted. And this is the most interesting and most useful degree of computational irreducibility. I will have more to say about this below.

Hiller and Isaacson's *Illiac Suite* (Lejaren A. Hiller and Isaacson 1957) is the first piece of what can unambiguously be called computer music, and it is an algorithmic composition assembled using a toolkit of stochastic note generators and music transformation functions, as detailed in their book *Experimental Music* (Hiller and Isaacson 1959). This can be called the *toolkit approach* to algorithmic composition. The composer programs a chain of generators, including random variables, and transformations, including serial transformations, to produce a chunk of music. The chunks can then be edited by hand. Multiple chunks can be assembled into a composition by hand. The toolkit approach lives on in contemporary software systems such as Open Music (C. A. Agon, Gerard Assayag, and Bresson 2008), Common Music (Taube 2023, 2021), and many others. This is to date the most successful and widely used method of algorithmic composition.

Algorithms in the Toolkit

The toolkit approach includes transformations and generators from many sources:

Traditional Music Theory Fugue and other canonical forms, scales, transpositions and rescalings of pitch and time.

"Atonal Theory" or "Set Theory" Transpose, invert, and reverse, and other group actions on tone rows or other pitch collections (Rahn 1991)

Neo-Riemannian Theory Mathematizes voice-leading and discovers groups acting on notes, scales, and chords, together with their symmetries (Tymoczko 2006a, 2006b; Callender, Quinn, and Tymoczko 2008).

Random Variables Borrowed from mathematics and used to generate series of musical properties such as pitches.

Dynamical Systems and Cellular Automata Borrowed from mathematics and used to generate series of musical properties such as pitches (Miranda 1993).

Fractals Borrowed from mathematics and used to generate some or all aspects of musical scores (Miranda 2001; Madden 2007). Some composers, such as myself, prefer to use an algorithm, such as a Lindenmayer system (Prusinkiewicz and Lindenmayer, 1996 [1991]; Prusinkiewicz 1986; Gogins 1992b) or iterated function system (IFS) (Barnsley, 1993 [1988][b]; Gogins 1991) that will generate an entire piece based on fractals or other mathematical methods, without need of further editing or assembling.

Evolutionary Computing Applying evolutionary computing to any of the above (Miranda and Biles 2007).

Intrinsic vs. Extrinsic Algorithms

Most algorithms in the toolkit are computationally irreducible. However, there is a critical difference between applying, *e.g.*, a series of transformations from the General Contextual Group (Fiore and Satyendra 2005), even if the transformations are selected at random, versus simply applying a random variable to scramble or transpose a chord. I will define as *intrinsic* all algorithms that are based upon primitive elements of music and follow musically acceptable transformations of the primitives. For example, any application of the GCG is going to generate a chord progression that sounds more or less musically well-formed, because the group instantiates an analysis of some aspects of neo-Riemannian theory. I will define as *extrinsic* all algorithms that are based on non-musical primitives, use transformations of the primitives that are not necessarily musically acceptable, and end up mapping the productions of the algorithms to musical scores.

In other words, while both procedures are computationally irreducible, the GCG procedure will generate many more musically acceptable results. That is because the GCG models the syntax of some aspects of musical practice. The elements of the GCG are chords, and the actions of the GCG implement some common transformations of one chord to another. Other such syntaxes include abstracting the pragmatic rules of voice-leading, abstracting the network of chord progressions by scale degree in functional harmony, and so on.

Note well, a fractal or other mathematically derived algorithm may be modified to be based on, or to transform, musical primitives.

Live Coding

The more recent method of algorithmic composition known as *live coding* can be considered a variant of the toolkit approach. A live coding system for music consists of a toolkit of routines that are assembled into a music-generating graph during a live performance by interpreting real-time commands in a domain-specific language. Live coding systems have tools and commands for both high-level representations of music (notes, loops,

chords, scales, musical transformations, *etc.*) and sound synthesis (oscillators, envelope generators, filters, *etc.*). An overview of the field can be gleaned from the TOPLAP web site (toplap.org 2023) and the *Oxford Handbook of Algorithmic Composition* (McLean and Dean 2018). I have some experience with TidalCycles (computer platform, Haskell implementation) (McLean et al. 2023) and Strudel (a JavaScript version of TidalCycles, Web browser platform, JavaScript implementation) (Roos et al. 2023).

Machine Learning

Recently it has become possible to compose music using large language models (LLMs) such as ChaptGPT. This can be called the *machine learning* approach to computer music. (I prefer the term *machine learning* to *artificial intelligence* because the software is not intelligent, but it is trainable and so in some sense it is learning). I discuss only LLMs as they are currently the most influential method of machine learning. Briefly, the approach is based on emulating biology, specifically, on simulating at a high level of abstraction the behavior of nerve cells. A neural network is built up from layers of simulated neurons that connect with each other; the connections have tunable weights that control the output behavior of neurons in one layer given inputs from connected neurons in other layers. First the data, a large corpus of texts, is syntactically broken up into tokens. "Attention heads" are trained (in parallel) to learn weights indicating how tokens within the entire data set are related syntactically and semantically; each head is trained to reflect some different aspect of the data. The trained heads are then used to transform the input (the prompt) *plus* current output (hence the term "*self*-attention") to a new input from which the next token of output is predicted (Vaswani et al. 2017). This attention mechanism, and other heuristic mechanisms, have been found to greatly increase the power of the neural networks. In particular, the attention mechanism makes it possible to train the network on a very large body of data more efficiently and without much human intervention. For more detail, see (Zhang et al. 2023) and OpenAI's paper on their current LLM architecture (OpenAI 2023a). For working examples of how LLMs can be used to compose music, see Jukebox (OpenAI 2023b), Gonsalves (Gonsalves 2021), and Ocampo et al. (Ocampo et al. 2023).

Although it is early days for machine learning, contemporary experience has led to a number of reviews and critical studies of the capabilities and limitations of machine learning. From the skeptical side, see (Dale 2021). For an amusing and instructive series of dialogues between all sides, see (Aaronson 2023). This experience makes it possible to identify a few important things about the mathematical foundations of algorithmic composition based on machine learning:

Computational opacity All agree that LLMs can generate amazing, even spooky, results without anyone understanding much about what is going on in the neural network. We have a perfect understanding of each component in the LLM, because these components are actually quite simple, but we have no idea at all how an LLM like ChatGPT can conduct a fact-filled conversation with one in perfect English. The details are scattered through a trillion or so neural network weights in the LLM.

Computational opacity is a form of computational irreducibility, but it goes far beyond irreducibility because, with computational opacity, we cannot obtain even a partial understanding of the actual computations performed by the software. We have taken one irreducible program (the untrained LLM) and used it to build another irreducible program (the trained LLM)! Even though an untrained LLM is computationally irreducible, we still have a perfect understanding of how it actually works; but it seems very likely that we will not, at least in practice, ever obtain even a partial understanding of how the *trained* LLM actually works.

Hallucination Refers to the tendency of LLMs to generate factually incorrect responses to prompts. It is a reminder that the software has no sense of reality and no means of comparing what it generates with the real world. I suspect that hallucinations arise from the human mistakes, conflicting goals, and outright lies represented in the training data. Ways of dealing with hallucinations are being investigated; for one approach, see (Christiano et al. 2017).

Unoriginality LLMs generate responses to prompts based on high-dimensional correlations that the LLMs have automatically discovered in the training data – data that *we* have provided. This is a self-referential situation. When we converse with ChatGPT we are looking at ourselves in a mirror – indeed, a fun-house mirror.

Artistic Results and Prospects

To date, not many pieces of algorithmically composed music have become popular or influential, even among composers and aficionados of art music and experimental music. Some pieces that have been influential are Iannis Xenakis' *La Légende D'Eer* (Xenakis 2005) and *Gendy 3* (Xenakis 1995), Charles Dodge's *Viola Elegy* (Dodge, April 1994), and Brian Eno's generative works (Eno 2023; Eno and Chilvers 2023). I have my own idiosyncratic list of different algorithmic composition systems, with my own choice of representative pieces (Gogins 2019).

The actual procedures followed by composers for algorithmic composition vary by genre, by composer, and by the software used. It is difficult to get a handle on the actual practices of any composers, let alone algorithmic composers. Composing can be a communal effort, as in much contemporary popular music, but art music is usually rather private, and algorithmically composed music even more so. However, IRCAM has published a series of books with chapters by composers on how they have used OpenMusic (C. Agon, Gérard Assayag, and Bresson 2006-2016, 2006a, 2006b, 2016). These are very useful. Profiles of composers in *Computer Music Journal* also can be useful. Here I will explain the general procedure that I myself follow.

I start with some kind of mathematical system that can be used to generate a set of musical notes, a score. The system needs to generate complex structure that can be changed by varying a relatively small number of numerical parameters or commands. It's often useful to select a recursive algorithm that, as the number of iterations approaches infinity, approaches a fixed point that is a fractal.

Such generative algorithms generally reflect processes in Nature that produce fractal-like forms, such as the patterns on seashells or the branching of plants. I have used chaotic dynamical systems, Lindenmayer systems, iterated function systems, and other systems.

Generally speaking, how to set the parameters in order to obtain a desired result is more or less opaque. This is well-known as the *inverse problem* (Graham and Demers 2021; Tu et al. 2023). But actually this is another form, once again, of computational irreducibility, meaning in this case that it is not intuitive how to infer the structure of an algorithm even after closely inspecting its results.

On the one hand this is a fault of the method; but on the other hand, and even more so, it is a virtue. In this way, and only in this way, can we generate scores that we would not otherwise be able to imagine. This, of course, is another fundamental motivation for pursuing algorithmic composition. And it's the most important motivation. *This kind of algorithmic composition actually amplifies our musical imagination.*

Now the question arises, how can such opaque methods be used to compose good music? It is difficult but by no means impossible, and here is the usual way I do it.

The parameters generally have a global effect on the final structure, that is, on the generated score. For example, an iterated function system consists of a number of affine transformations that are repeatedly applied to a set. Changing just one element of one transformation matrix can easily change every note in the generated score.

So, I pick one parameter and change its value, then listen to the result. I change the value again, and listen to the second result. Then, I choose the value that I prefer. I make more changes and listen again. Eventually I will find a value for that parameter that is more or less optimal – a sweet spot in the parameter space.

Then I pick another parameter and change its value in the same way, until I have a second sweet spot. During this process, the effect of the first parameter will probably change, so I go back to the first parameter and search again for the "sweet spot."

This can be repeated for any number of parameters, but it is a tedious process and does not make sense for more than a few parameters.

This procedure amounts to a sort of binary search through a set of possible parameter values so vast – indeed infinite – that a linear search is simply out of the question. But a binary search is far more efficient than a linear search. Furthermore, finding two or three "sweet spots" in a small set of controlling parameters – each of which has global effects on the entire score – can produce a surprisingly large improvement in the musicality of the result.

I see here an analogy with the way in which LLMs work. There are repeated searches in a parameter space equipped with a fitness function (as with the attention mechanism) at increasing levels of refinement (as with gradient descent).

Few algorithmic composers simply "hear music in their heads" and write software to render it, although that certainly happens. Most fool around producing various experimental chunks of music, refine them more or less as I have described, and assemble some into a finished composition.

Before I proceed to look at this kind of production from a mathematical point of view, I will summarize what I have learned about the mathematical foundations:

Uncomputability The set of musical compositions that are possible in the abstract (assuming that some pieces either last an infinitely long time, or that between any two pieces is a continuous path consisting of variations between the pieces) is *recursively uncomputable.*

Universality In spite of the uncomputability of compositions, they are *recursively enumerable*, so it is possible to approximate any possible composition as closely as one likes.

Irreducibility Compositional algorithms that have a strong analogy to Julia sets are *computationally irreducible.*

Intrinsic vs. Extrinsic Algorithms based on musical primitives and transformations do not produce nearly as many musically unacceptable results.

Opacity Compositional algorithms based on machine learning are not only computationally irreducible, but also computationally opaque in that we have essentially no insight into the meaning of the steps followed by the LLM.

Mappability Compositional algorithms are nevertheless always mappable. This ultimately is because sets of compositions can be ordered in some way by *musical* criteria, while their corresponding generating parameters can be ordered lexically or numerically.

Intractability Producing a useful parametric map of some subset of compositions is infeasibly compute-intensive.

Hallucination LLMs that are supposed to provide true or useful outputs sometimes just make stuff up. But this means that material generated by an LLM in response to a prompt cannot be trusted to be true or useful. It is necessary for a person, indeed an expert, to evaluate the material. It is by no means clear at this time whether an expert equipped with a LLM is more productive for creative work than that same expert without the LLM.

Unoriginality LLMs work by discovering high-dimensional correlations in large bodies of training data. This means that LLMs can select, summarize, and vary — but they cannot generate an output that is not correlated with the training data. In other words, there is a limit to their originality. However, it is by no means clear at this time whether that limit is well below, or well above, the creativity of experts in the field from which the training data was drawn.

I will now put forward some conjectures based on these foundations.

Limitations

At this time and for the foreseeable future, no form of artificial intelligence is conscious or has its own goals. Therefore, for the foreseeable future, human composers must and

will play a irreplaceable role in algorithmic composition. This involves selecting a subset of possible compositions to study, evaluating the musical quality of each composition in the subset, and varying the parameters or prompts that generate the pieces. This follows from hallucination and unoriginality.

Incomputability, irreducibility, and opacity set objective limits on how much understanding composers can gain into the working of their algorithms and of the music generated by them. This is both a limitation and an advantage. In practice, it is not possible to determine in advance just where those limits lie.

Sophisticated forms of algorithmic composition are compute-intensive, and can be computationally intractable.

Prospects

Computer power will continue to increase. This will most likely make algorithmic composition both more productive and more important.

There is a similarity between a composer's experience with a toolkit of algorithms, the transformation of prompts into responses by an LLM, and exploring the parametric map of a fractal compositional algorithm. In all cases, starting with an initial trial, a final composition is approached by a descending, zigzag search through a space representing musical possibilities of differing value, until the search comes to rest in some local optimum. This search can be greatly speeded up by using intrinsically musical algorithms that tend not to produce as many musically unacceptable results. For example, rather than representing scores as notes on piano rolls, e.g. planes or cubes, one can represent scores as more or less fleeting chord progressions in chord spaces (Gogins 2006, In preparation).

Every method that speeds up the search process will make algorithmic composition more productive. In particular, the growth of live coding demonstrates that the toolkit approach to algorithmic composition has a future. The underlying reason is that live coding supports faster searching, due to concise commands and immediately audible feedback. Spending time doing live coding also increases the composer's insight into the tools.

In Sum

The main result here is that the major approaches to algorithmic composition — trial and error with a toolkit of algorithms, live coding, exploration of fractals, and machine learning — share this fundamental business of zigzagging down a slope on a landscape of evaluations to rest in a local optimum. This result is proved by the simple fact that the generated music and/or the parameters used to generate it can be ordered. The dimensionality of the parameter space is secondary, as it can be reduced to one or two dimensions by means of a Hilbert index. Note that searching for solutions or optimizations in many domains is known to be computationally expensive.

Future developments in artificial intelligence may have a significant impact on algorithmic composition. For example, machine learning has been applied to solving the

inverse problem for discovering the parameters of fractal algorithms (Tu et al. 2023). It might then be possible to represent an existing score, or scores, as fractal parameters and then work with these parameters to vary or interpolate between such pieces. This does not overcome computational irreducibility, as it substitutes the opacity of machine learning for the irreducibility of the inverse problem, yet it still might be very useful.

Algorithmic compositions based on current LLMs are easier to produce, but will not usually be musically original; while algorithmic compositions based on toolkits, live coding, or fractals can be musically original, but are inherently more difficult to produce.

As for algorithmic compositions based on fractals, sometimes an analytical understanding of the mathematics can be used as a guide to composition, but this tends in my experience to be of limited use. More often, the only way to penetrate the fog of incomputability, irreducibility, opacity, and intractability is to explore the geometrical order in a subset of compositions. This can be done either by trial and error, or by literally plotting a map of the subset of compositions. One might say that with trial and error one plots a sparse map of fully defined features in a territory, and with a parametric map one explores a densely mapped territory with partially defined features.

Progress in algorithmic composition seems likely to depend on speeding up the composer's workflow, whether in parametric composition, in live coding, in algorithmic composition toolkits, or in machine learning; and, in the long run, on defining more musically compact and intelligible spaces of musical possibility, based on a wider variety of intrinsically musical algorithms. This in my view is one of the most important tasks of future algorithmic composers: identify and refine such processes.

However, it would be unwise to exclude extrinsic algorithms from the toolkit, as they may expand one's hearing of what is musically acceptable.

Finally, if P = NP is ever proved, someone may succeed in creating an algorithm that can compute a dense parametric map of a compositionally universal algorithm in polynomial time. The result might be, again, a God's-eye view of possible structure in music. But I'm not holding my breath, and I think we must continue to work within a garden, it is to be hoped one that grows ever larger and richer, of disparate compositional algorithms.

References

Aaronson, Scott. 2005. Np-complete problems and physical reality. *ACM SIGACT News* 36 (1): 30–52.

———. 2023. *Shtetl-optimized: the blog of scott aaronson.* Available online at: `https://scottaaronson.blog/`. Accessed 10 May 2023.

Agon, Carlos, Gérard Assayag, and Jean Bresson. 2006a. *The om composer's book 1.* Vol. 2006. Editions Delatour France/Ircam-Centre Pompidou.

———. 2006b. The om composer's book 2. *Editions Delatour/IRCAM Centre Pompidou* 2008.

Agon, Carlos, Gérard Assayag, and Jean Bresson. 2016. The om composer's book 3. *Editions Delatour/IRCAM Centre Pompidou* 2016.

——. 2006-2016. *The om composers book.* Available online at: `http://repmus.ircam.fr/openmusic/ombook`. Accessed 11 May 2023.

Agon, Carlos Augusto, Gerard Assayag, and Jean Bresson. 2008. *Openmusic.* `Http://recherche.ircam.fr/equipes/repmus/OpenMusic/`.

Ariza, Christopher. 2023. *Algorithmic.net.* Available online at: `http://algorithmic.net/`. Accessed 13 May 2023.

Arora, Sanjeev, and Boaz Barak. 2009. *Computational complexity: a modern approach.* Cambridge University Press. ISBN: 9780521424264. `https://theory.cs.princeton.edu/complexity/book.pdf`.

Barnsley, Michael F, and Stephen Demko. 1985. Iterated function systems and the global construction of fractals. *Proceedings of the Royal Society of London. A. Mathematical and Physical Sciences* 399 (1817): 243–275.

Barnsley, Michael F, and Alan D Sloan. 1989. Fractal image compression. In *Proc. scientific data compression workshop,* 351–365.

Barnsley, Michael F. 1993 [1988](a). *Fractals everywhere.* 2nd. Boston: Academic Press Professional.

——. 1993 [1988](b). *Fractals everywhere.* 2nd. Boston: Academic Press Professional.

Blum, Lenore, and Steve Smale. 1993. The gödel incompleteness theorem and decidability over a ring. In *From topology to computation: proceedings of the smalefest,* 321–339. Springer.

Braverman, Mark, and Michael Yampolsky. 2006. Non-computable julia sets. *Journal of the American Mathematical Society* 19 (3): 551–578.

——. 2009. *Computability of julia sets.* Vol. 23. Available online at: `https://arxiv.org/pdf/math/0610340.pdf`. Accessed 11 May 2023). Springer.

Callender, Clifton, Ian Quinn, and Dmitri Tymoczko. 2008. Generalized voice-leading spaces. *Science* 320 (5874): 346–348.

Christiano, Paul F, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. 2017. Deep reinforcement learning from human preferences. *Advances in neural information processing systems* 30.

Copeland, B. Jack. 2024. The Church-Turing Thesis. In *The Stanford encyclopedia of philosophy,* Winter 2024, edited by Edward N. Zalta and Uri Nodelman. Metaphysics Research Lab, Stanford University.

Dale, Robert. 2021. Gpt-3: what's it good for? *Natural Language Engineering* 27 (1): 113–118.

Davis, Martin. 2018. *The universal computer: the road from leibniz to turing.* 3rd. Boca Raton, FL: CRC Press. ISBN: 978-1138502086.

Dodge, Charles. April 1994. Any Resemblance is Purely Coincidental. Chap. Viola Elegy. NA043. New Albion Records.

Eno, Brian. 2023. *Generative music.* Available online at: `https://inmotionmagazine.com/eno1.html`. Accessed 10 May 2023.

Eno, Brian, and Peter Chilvers. 2023. *Bloom: 10 worlds by brian eno and peter chilvers. world 1: origin.* Available online at: `https://music.youtube.com/watch?v=rGAUQmfV6w4`. Accessed 10 May 2023.

Fernández, Jose D, and Francisco Vico. 2013. Ai methods in algorithmic composition: a comprehensive survey. *Journal of Artificial Intelligence Research* 48:513–582.

Fiore, T.M., and R. Satyendra. 2005. Generalized Contextual Groups. *Music Theory Online* 11 (3).

Glass, Philip. 2015. *Words without music: a memoir.* WW Norton & Company.

Gödel, Kurt. 1986. On formally undecidable propositions of *Principia Mathematica* and related systems i. In *Kurt gödel: collected works, volume i: publications 1929–1936,* edited by Solomon Feferman, John W. Dawson Jr., Stephen C. Kleene, Gregory H. Moore, Robert M. Solovay, and Jean van Heijenoort, 144–195. Originally published in German in 1931. Oxford University Press. ISBN: 9780195039641.

Gogins, Michael. 1991. Iterated Function Systems Music. *Computer Music Journal* 15 (1): 34–42.

———. 1992a. ...how i became obsessed with finding a mandelbrot set for sounds. *News of Music* 13:129–139.

———. 1992b. Fractal Music with String Rewriting Grammars. *News of Music* 13:146–170.

———. 2006. Score generation in voice-leading and chord spaces. In *Proceedings of the 2006 international computer music conference,* edited by Georg Essl and Ichiro Fujinaga. San Francisco, California: International Computer Music Association.

———. 2019. *More rant-like musings on algorithmic composition software.* Available online at: `https://michaelgogins.tumblr.com/post/188345210028/algorithmiccompositionsystems`. Accessed 9 May 2023.

———. In preparation. Parametric composition of score graphs.

Gonsalves, Robert A. 2021. *Ai-tunes: creating new songs with artificial intelligence.* Available online at: `https://towardsdatascience.com/ai-tunes-creating-new-songs-with-artificial-intelligence-4fb383218146`. Accessed 9 May 2023.

Graham, Liam, and Matthew Demers. 2021. Applying neural networks to a fractal inverse problem. In *Recent developments in mathematical, statistical and computational sciences: the v ammcs international conference, waterloo, canada, august 18–23, 2019,* 157–165. Springer.

Hertling, Peter. 2005. Is the mandelbrot set computable? *Mathematical Logic Quarterly* 51 (1): 5–18. `https://doi.org/10.1002/malq.200310124`.

Hiller, Lejaren, and L.M. Isaacson, eds. 1959. *Experimental music: composition with an electronic computer.* New York, New York: McGraw–Hill.

Huffman, Carl. 2024. Pythagoras. In *The Stanford encyclopedia of philosophy,* Spring 2024, edited by Edward N. Zalta and Uri Nodelman. Metaphysics Research Lab, Stanford University.

Huffman, Carl A. 2014. *A history of pythagoreanism.* Cambridge University Press.

humdrum.org. 2023. *Musikalisches würfelspiel.* Available online at: `https://dice.humdrum.org/`. Accessed 11 May 2023.

Kawahira, Tomoki, and Masashi Kisaka. 2018. Julia sets appear quasiconformally in the mandelbrot set. *arXiv preprint arXiv:1804.00176.*

Lei, TAN. 1990. Similarity between the mandelbrot set and julia sets. *Commun. Math. Phys* 134:587–617.

Lejaren A. Hiller, Jr., and Leonard M. Isaacson. 1957. *Illiac suite for string quartet.* New York: New Music Editions.

Madden, Charles. 2007. *Fractals in music: introductory mathematics for musical analysis.* 2nd, Revised and Expanded. Salt Lake City, UT: High Art Press. ISBN: 9780967172774.

Mandelbrot, Benoit B. 1982. *The fractal geometry of nature.* W. H. Freeman, August. ISBN: 0716711869. `http://www.amazon.ca/exec/obidos/redirect?tag=citeulike09-20%5C&path=ASIN/071%206711869`.

McLean, A., and R.T. Dean. 2018. *The oxford handbook of algorithmic music.* Oxford Handbooks. Oxford University Press. ISBN: 9780190227005. `https://books.google.com/books?id=7XBGDwAAQBAJ`.

McLean, Alex, et al. 2023. *Tidal cycles.* Available online at: `https://tidalcycles.org/`. Accessed 14 May 2023.

Miranda, Eduardo Reck. 1993. Cellular automata music: an interdisciplinary project. *Interface* 22 (1): 3–21. `https://doi.org/10.1080/09298219308570616`.

———. 2001. *Composing music with computers* [in English]. Oxford; Boston: Focal Press. ISBN: 9780240515670.

Miranda, Eduardo Reck, and John Al Biles, eds. 2007. *Evolutionary computer music.* London: Springer. ISBN: 978-1-84628-598-8. `https://doi.org/10.1007/978-1-84628-599-5`.

Ocampo, Rodolfo, Josh Andres, Adrian Schmidt, Caroline Pegram, Justin Shave, Charlton Hill, Brendan Wright, and Oliver Bown. 2023. Using gpt-3 to achieve semantically relevant data sonificiation for an art installation. In *Artificial intelligence in music, sound, art and design: 12th international conference, evomusart 2023, held as part of evostar 2023, brno, czech republic, april 12–14, 2023, proceedings,* 212–227. Springer.

OpenAI. 2023a. *Gpt-4 technical report.* arXiv: 2303.08774 [cs.CL].

———. 2023b. *Jukebox.* Available online at: `https://openai.com/research/jukebox/`. Accessed 10 May 2023.

Patrick, Edward A, Douglas R Anderson, and Friend K Bechtel. 1968. Mapping multi-dimensional space to one dimension for computer output display. In *Proceedings of the 1968 23rd acm national conference,* 511–515.

Peitgen, Heinz-Otto, Hartmut Jürgens, Dietmar Saupe, Heinz-Otto Peitgen, Hartmut Jürgens, and Dietmar Saupe. 2004. The mandelbrot set: ordering the julia sets. *Chaos and Fractals: New Frontiers of Science,* 783–837.

Potter, Keith. 2002. *Four musical minimalists: la monte young, terry riley, steve reich, philip glass.* Vol. 11. Cambridge University Press.

Prusinkiewicz, P. 1986. Score generation with L-systems. In *Proceedings of the 1986 international computer music conference,* 455–457.

Prusinkiewicz, Przemyslaw, and Artistid Lindenmayer. 1996 [1991]. *The algorithmic beauty of plants.* Available online at `http://algorithmicbotany.org/papers`. New York: Spring-Verlag.

Rahn, John. 1991. *Basic atonal theory.* 2nd. New York: Longman. ISBN: 9780028721453.

Reich, Steve. 2023. *Music as a gradual process.* Available online at: `https://static1.squarespace.com/static/50e79ec7e4b07dba60068e4d/t/515707b0e4b0ec1768d61b17/1364658096256/Reich.pdf`. Accessed 11 May 2023.

Roos, Felix, et al. 2023. *Strudel.* Available online at: `https://strudel.tidalcycles.org/`. Accessed 14 May 2023.

Schwarz, K. Robert. 1980. Steve reich: music as a gradual process: part i. *Perspectives of New Music* 19 (1/2): 373–392. ISSN: 00316016, accessed May 11, 2023. `http://www.jstor.org/stable/832600`.

Taube, Rick. 2021. *Musx.* Available online at: `https://github.com/musx-admin/musx`. Accessed 9 May 2023.

———. 2023. *Common Music.* `Https://commonmusic.sourceforge.net/`.

Tennant, Neil. 2023. Logicism and Neologicism. In *The Stanford encyclopedia of philosophy,* Winter 2023, edited by Edward N. Zalta and Uri Nodelman. Metaphysics Research Lab, Stanford University.

toplap.org. 2023. *Toplap.* Available online at: `https://toplap.org/about/`. Accessed 14 May 2023.

Tu, Cheng-Hao, Hong-You Chen, David Carlyn, and Wei-Lun Chao. 2023. Learning fractals by gradient descent. *arXiv preprint arXiv:2303.12722.*

Tymoczko, Dmitri. 2006a. The Geometry of Musical Chords. *Science* 313:72–74.

———. 2006b. The geometry of musical chords. *Science* 313 (5783): 72–74.

Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems* 30.

Wolfram, Stephen. 1985. Undecidability and intractability in theoretical physics. *Physical Review Letters* 54 (8): 735.

Xenakis, Iannis. 1995. Xenakis: Ais, Gendy3, Taurhiphanie, Thallein / Nee, Steiger. Chap. Gendy3. 45086. Neuma.

———. 2005. *La légende d'eer.* Edited by Gerard Pape and Iannis Xenakis. Xenakis, Iannis, 1922-2001. Electronic works ; Program notes in English, French, and German ([21] p. : ill.) in container. New York, NY: Mode.

Zach, Richard. 2023. Hilbert's Program. In *The Stanford encyclopedia of philosophy,* Spring 2023, edited by Edward N. Zalta and Uri Nodelman. Metaphysics Research Lab, Stanford University.

Zhang, Chaoning, Chenshuang Zhang, Sheng Zheng, Yu Qiao, Chenghao Li, Mengchun Zhang, Sumit Kumar Dam, Chu Myaet Thwal, Ye Lin Tun, Le Luang Huy, et al. 2023. A complete survey on generative ai (aigc): is chatgpt from gpt-4 to gpt-5 all you need? *arXiv preprint arXiv:2303.11717.*

Zwirn, Herve. 2015. Computational irreducibility and computational analogy. *HAL* 2015.